

```

[acadgild@localhost ~]$ pwd
/home/acadgild
[acadgild@localhost ~]$ ls
Desktop    eclipse-jee-neon-M3-linux-gtk-x86_64.tar.gz  Pictures  workspace
Documents  hadoop                                         Public
Downloads  metastore_db                                  Templates
eclipse    Music                                          Videos
[acadgild@localhost ~]$ cd Downloads
[acadgild@localhost Downloads]$ ls
apache-hive-0.14.0-bin.tar.gz.part  wordcount.zip
hadoop-mapreduce-examples-2.9.0.jar
[acadgild@localhost Downloads]$ cp /hadoop-mapreduce-examples-2.9.0.jar /home/acadgild/Desktop
cp: cannot stat '/hadoop-mapreduce-examples-2.9.0.jar': No such file or directory
[acadgild@localhost Downloads]$ cp hadoop-mapreduce-examples-2.9.0.jar /home/acadgild/Desktop
[acadgild@localhost Downloads]$ ls
apache-hive-0.14.0-bin.tar.gz.part  wordcount.zip
hadoop-mapreduce-examples-2.9.0.jar
[acadgild@localhost Downloads]$ cp wordcount.zip /home/acadgild/Desktop
[acadgild@localhost Downloads]$ pwd
/home/acadgild/Downloads
[acadgild@localhost Downloads]$ cd /home/acadgild/Desktop
[acadgild@localhost Desktop]$ ls
gnome-terminal.desktop  new file
hadoop-mapreduce-examples-2.9.0.jar  wordcount.zip
[acadgild@localhost Desktop]$ clear

```

```

[acadgild@localhost Desktop]$ cat file1.txt
cat: file1.txt: No such file or directory
[acadgild@localhost Desktop]$ cat > file1.txt
A B C
C A B
B B A
B C A
A C C
^C

```

```

[acadgild@localhost Desktop]$ start-all.sh
This script is Deprecated. Instead use start-dfs.sh and start-yarn.sh
18/03/13 14:53:11 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Starting namenodes on [localhost]
localhost: starting namenode, logging to /usr/local/hadoop-2.6.0/logs/hadoop-acadgild-namenode-localhost.localdomain.out
localhost: starting datanode, logging to /usr/local/hadoop-2.6.0/logs/hadoop-acadgild-datanode-localhost.localdomain.out
Starting secondary namenodes [0.0.0.0]
0.0.0.0: starting secondarynamenode, logging to /usr/local/hadoop-2.6.0/logs/hadoop-acadgild-secondarynamenode-localhost.localdomain.out
18/03/13 14:53:30 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
starting yarn daemons
starting resourcemanager, logging to /usr/local/hadoop-2.6.0/logs/yarn-acadgild-resourcemanager-localhost.localdomain.out
localhost: starting nodemanager, logging to /usr/local/hadoop-2.6.0/logs/yarn-acadgild-nodemanager-localhost.localdomain.out
[acadgild@localhost Desktop]$ clear

```

```
[acadgild@localhost Desktop]$ ls
file1.txt          hadoop-mapreduce-examples-2.9.0.jar  wordcount.zip
gnome-terminal.desktop  new file
[acadgild@localhost Desktop]$ hadoop fs -put file1.txt
18/03/13 14:53:56 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
[acadgild@localhost Desktop]$ hadoop fs put file1.txt
put: Unknown command
Did you mean -put? This command begins with a dash.
[acadgild@localhost Desktop]$ hadoop fs -ls
18/03/13 14:54:33 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 3 items
-rw-r--r--    1 acadgild supergroup          31 2018-03-13 14:53 file1.txt
drwxr-xr-x   - acadgild supergroup           0 2015-11-17 02:03 oozie-acad
drwxr-xr-x   - acadgild supergroup           0 2015-11-17 02:00 share
```

```
[acadgild@localhost Desktop]$ ls
file1.txt          hadoop-mapreduce-examples-2.9.0.jar  wordcount.zip
gnome-terminal.desktop  new file
[acadgild@localhost Desktop]$ hadoop fs -put file1.txt
18/03/13 14:54:57 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
put: `file1.txt': File exists
[acadgild@localhost Desktop]$ hadoop fs -cat file1.txt
18/03/13 14:55:16 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
A B C
C A B
B B A
B C A
A C C
```

## Word Median

```
[acadgild@localhost Desktop]$ hadoop jar hadoop-mapreduce-examples-2.9.0.jar wordmedian file1.txt wmedian
18/03/13 15:10:41 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
18/03/13 15:10:41 INFO client.RMPProxy: Connecting to ResourceManager at /0.0.0.0:8032
18/03/13 15:10:42 WARN mapreduce.JobSubmitter: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
18/03/13 15:10:43 INFO input.FileInputFormat: Total input paths to process : 1
18/03/13 15:10:43 INFO mapreduce.JobSubmitter: number of splits:1
18/03/13 15:10:43 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1520933014970_0002
18/03/13 15:10:43 INFO impl.YarnClientImpl: Submitted application application_1520933014970_0002
18/03/13 15:10:43 INFO mapreduce.Job: The url to track the job: http://localhost:8088/proxy/application_1520933014970_0002/
18/03/13 15:10:43 INFO mapreduce.Job: Running job: job_1520933014970_0002
18/03/13 15:10:51 INFO mapreduce.Job: Job job_1520933014970_0002 running in uber mode : false
18/03/13 15:10:51 INFO mapreduce.Job: map 0% reduce 0%
18/03/13 15:10:58 INFO mapreduce.Job: map 100% reduce 0%
18/03/13 15:11:05 INFO mapreduce.Job: map 100% reduce 100%
18/03/13 15:11:05 INFO mapreduce.Job: Job job_1520933014970_0002 completed successfully
18/03/13 15:11:05 INFO mapreduce.Job: Counters: 49
```

### File System Counters

```
FILE: Number of bytes read=16
FILE: Number of bytes written=211897
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=141
HDFS: Number of bytes written=5
HDFS: Number of read operations=6
HDFS: Number of large read operations=0
HDFS: Number of write operations=2
```

### Job Counters

```
Launched map tasks=1
Launched reduce tasks=1
Data-local map tasks=1
Total time spent by all maps in occupied slots (ms)=3888
Total time spent by all reduces in occupied slots (ms)=3962
Total time spent by all map tasks (ms)=3888
Total time spent by all reduce tasks (ms)=3962
Total vcore-seconds taken by all map tasks=3888
Total vcore-seconds taken by all reduce tasks=3962
Total megabyte-seconds taken by all map tasks=3981312
Total megabyte-seconds taken by all reduce tasks=4057088
```

#### Map-Reduce Framework

```
Map input records=5
Map output records=15
Map output bytes=120
Map output materialized bytes=16
Input split bytes=110
Combine input records=15
Combine output records=1
Reduce input groups=1
Reduce shuffle bytes=16
Reduce input records=1
Reduce output records=1
Spilled Records=2
Shuffled Maps =1
Failed Shuffles=0
Merged Map outputs=1
GC time elapsed (ms)=125
CPU time spent (ms)=1060
Physical memory (bytes) snapshot=294162432
Virtual memory (bytes) snapshot=4113768448
Total committed heap usage (bytes)=165810176
```

#### Shuffle Errors

```
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0
```

#### File Input Format Counters

```
Bytes Read=31
```

#### File Output Format Counters

```
Bytes Written=5
```

```
The median is: 1
```

```
[acadgild@localhost Desktop]$ hadoop fs -ls wmedian
18/03/13 15:14:28 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 2 items
-rw-r--r--  1 acadgild supergroup          0 2018-03-13 15:11 wmedian/_SUCCESS
-rw-r--r--  1 acadgild supergroup          5 2018-03-13 15:11 wmedian/part-r-000000
```

```
[acadgild@localhost Desktop]$ hadoop fs -cat wmedian/part-r-000000
18/03/13 15:15:27 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
1      15
[acadgild@localhost Desktop]$
```

## Wordmean

```
[acadgild@localhost Desktop]$ hadoop jar hadoop-mapreduce-examples-2.9.0.jar wordcount file1.txt wmean
18/03/13 15:16:46 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
18/03/13 15:16:47 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032
18/03/13 15:16:48 INFO input.FileInputFormat: Total input paths to process : 1
18/03/13 15:16:48 INFO mapreduce.JobSubmitter: number of splits:1
18/03/13 15:16:48 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1520933014970_0003
18/03/13 15:16:49 INFO impl.YarnClientImpl: Submitted application application_1520933014970_0003
18/03/13 15:16:49 INFO mapreduce.Job: The url to track the job: http://localhost:8088/proxy/application_1520933014970_0003/
18/03/13 15:16:49 INFO mapreduce.Job: Running job: job_1520933014970_0003
18/03/13 15:16:57 INFO mapreduce.Job: Job job_1520933014970_0003 running in uber mode : false
18/03/13 15:16:57 INFO mapreduce.Job: map 0% reduce 0%
18/03/13 15:17:03 INFO mapreduce.Job: map 100% reduce 0%
18/03/13 15:17:10 INFO mapreduce.Job: map 100% reduce 100%
18/03/13 15:17:10 INFO mapreduce.Job: Job job_1520933014970_0003 completed successfully
```

```
[acadgild@localhost Desktop]$ hadoop fs -ls wmean
18/03/13 15:23:00 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 2 items
-rw-r--r-- 1 acadgild supergroup 0 2018-03-13 15:17 wmean/_SUCCESS
-rw-r--r-- 1 acadgild supergroup 19 2018-03-13 15:17 wmean/part-r-000000
0
```

```
[acadgild@localhost Desktop]$ hadoop fs -cat wmean/part-r-000000
18/03/13 15:23:45 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
count 15
length 15
_
```

```
18/03/13 15:17:10 INFO mapreduce.Job: Counters: 49
File System Counters
  FILE: Number of bytes read=39
  FILE: Number of bytes written=212203
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=141
  HDFS: Number of bytes written=19
  HDFS: Number of read operations=6
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=2

Job Counters
  Launched map tasks=1
  Launched reduce tasks=1
  Data-local map tasks=1
  Total time spent by all maps in occupied slots (ms)=4032
  Total time spent by all reduces in occupied slots (ms)=3892
  Total time spent by all map tasks (ms)=4032
  Total time spent by all reduce tasks (ms)=3892
  Total vcore-seconds taken by all map tasks=4032
  Total vcore-seconds taken by all reduce tasks=3892
  Total megabyte-seconds taken by all map tasks=4128768
  Total megabyte-seconds taken by all reduce tasks=3985408
```

#### Map-Reduce Framework

Map input records=5  
Map output records=30  
Map output bytes=435  
Map output materialized bytes=39  
Input split bytes=110  
Combine input records=30  
Combine output records=2  
Reduce input groups=2  
Reduce shuffle bytes=39  
Reduce input records=2  
Reduce output records=2  
Spilled Records=4  
Shuffled Maps =1  
Failed Shuffles=0  
Merged Map outputs=1  
GC time elapsed (ms)=125  
CPU time spent (ms)=1050  
Physical memory (bytes) snapshot=296333312  
Virtual memory (bytes) snapshot=4113768448  
Total committed heap usage (bytes)=165810176

#### Shuffle Errors

BAD\_ID=0  
CONNECTION=0  
IO\_ERROR=0  
WRONG\_LENGTH=0  
WRONG\_MAP=0  
WRONG\_REDUCE=0

#### File Input Format Counters

Bytes Read=31

#### File Output Format Counters

Bytes Written=19

The mean is: 1.0

```
[acadgild@localhost Desktop]$ hadoop fs -cat wmean/part-r-00000
18/03/13 15:23:45 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
count    15
length   15
```

## Word Standard Deviation

```
[acadgild@localhost Desktop]$ hadoop jar hadoop-mapreduce-examples-2.9.0.jar wordstandarddeviation file1.txt wstd
18/03/13 15:26:05 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
18/03/13 15:26:06 INFO client.RMPProxy: Connecting to ResourceManager at /0.0.0.0:8032
18/03/13 15:26:07 INFO input.FileInputFormat: Total input paths to process : 1
18/03/13 15:26:07 INFO mapreduce.JobSubmitter: number of splits:1
18/03/13 15:26:07 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1520933014970_0004
18/03/13 15:26:07 INFO impl.YarnClientImpl: Submitted application application_1520933014970_0004
18/03/13 15:26:07 INFO mapreduce.Job: The url to track the job: http://localhost:8088/proxy/application_1520933014970_0004/
18/03/13 15:26:07 INFO mapreduce.Job: Running job: job_1520933014970_0004
18/03/13 15:26:15 INFO mapreduce.Job: Job job_1520933014970_0004 running in uber mode : false
18/03/13 15:26:15 INFO mapreduce.Job: map 0% reduce 0%
18/03/13 15:26:21 INFO mapreduce.Job: map 100% reduce 0%
18/03/13 15:26:27 INFO mapreduce.Job: map 100% reduce 100%
18/03/13 15:26:27 INFO mapreduce.Job: Job job_1520933014970_0004 completed successfully
```

```
18/03/13 15:26:27 INFO mapreduce.Job: Counters: 49
```

#### File System Counters

```
FILE: Number of bytes read=56
FILE: Number of bytes written=212395
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=141
HDFS: Number of bytes written=29
HDFS: Number of read operations=6
HDFS: Number of large read operations=0
HDFS: Number of write operations=2
```

#### Job Counters

```
Launched map tasks=1
Launched reduce tasks=1
Data-local map tasks=1
Total time spent by all maps in occupied slots (ms)=3610
Total time spent by all reduces in occupied slots (ms)=3883
Total time spent by all map tasks (ms)=3610
Total time spent by all reduce tasks (ms)=3883
Total vcore-seconds taken by all map tasks=3610
Total vcore-seconds taken by all reduce tasks=3883
Total megabyte-seconds taken by all map tasks=3696640
Total megabyte-seconds taken by all reduce tasks=3976192
```

#### Map-Reduce Framework

```
Map input records=5
Map output records=45
Map output bytes=660
Map output materialized bytes=56
Input split bytes=110
Combine input records=45
Combine output records=3
Reduce input groups=3
Reduce shuffle bytes=56
Reduce input records=3
Reduce output records=3
Spilled Records=6
Shuffled Maps =1
Failed Shuffles=0
Merged Map outputs=1
GC time elapsed (ms)=137
CPU time spent (ms)=950
Physical memory (bytes) snapshot=295899136
Virtual memory (bytes) snapshot=4113764352
Total committed heap usage (bytes)=165810176
```

#### Shuffle Errors

```
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0
```

#### File Input Format Counters

```
Bytes Read=31
```

#### File Output Format Counters

```
Bytes Written=29
```

The standard deviation is: 0.0

```
[acadgild@localhost Desktop]$ hadoop fs -ls wstd
18/03/13 15:31:08 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 2 items
-rw-r--r--  1 acadgild supergroup          0 2018-03-13 15:26 wstd/_SUCCESS
-rw-r--r--  1 acadgild supergroup        29 2018-03-13 15:26 wstd/part-r-000000
[acadgild@localhost Desktop]$ hadoop fs -cat wstd/part-r-000000
18/03/13 15:31:34 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
count    15
length   15
square   15
[acadgild@localhost Desktop]$
```